

Intro__to__R__2016__Code.R

breannechryst

Wed Feb 10 16:38:18 2016

```
## Introduction to R
## Breanne Chryst and Sara Bastomski
## CSSSI StatLab
## Feb 19, 2016

#####
## R Basics
#####

## R can be used as a calculator, it works as expected:
2+3

## [1] 5

exp(2)

## [1] 7.389056

5^(2)

## [1] 25

3*2+4

## [1] 10

(2*4+1)/2

## [1] 4.5

## The print command:
print("Hello World")

## [1] "Hello World"

print(5*2-11)

## [1] -1

## Assigning a variable
x <- 5 # 5 has now been assigned to the variable x

# the "<-" assigns the value on the right to the name on the left. Made by "alt -"

x
```

```
## [1] 5
```

```
x^2
```

```
## [1] 25
```

```
## Creating a vector:
```

```
y <- c(3,7,5,1,2,3,2,5,5) # "c()" concatenates, creating a vector
```

```
## Extracting values of a vector:
```

```
y[2]
```

```
## [1] 7
```

```
3:5 # the whole numbers from 3 to 5
```

```
## [1] 3 4 5
```

```
y[3:5] # extracts the 3rd to 5th elements in the vector y
```

```
## [1] 5 1 2
```

```
sub.y <- y[3:5]
```

```
sub.y
```

```
## [1] 5 1 2
```

```
## Other ways to make vectors:
```

```
array(data = 0, dim = 3)
```

```
## [1] 0 0 0
```

```
seq(from = 1, to = 4, by = 1)
```

```
## [1] 1 2 3 4
```

```
seq(1,4,0.5)
```

```
## [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0
```

```
rep(x = "cat", times = 3)
```

```
## [1] "cat" "cat" "cat"
```

```
rep(c("cat", 4, x^(2)), each = 2)
```

```
## [1] "cat" "cat" "4" "4" "25" "25"
```

```
rep(c("cat", 4, x^(2)), times = 2)
```

```
## [1] "cat" "4" "25" "cat" "4" "25"
```

```
## "matrix()" creates a matrix from the values entered:
```

```
z <- matrix(y, nrow=3) # This is filled by column  
z
```

```
##      [,1] [,2] [,3]  
## [1,]    3    1    2  
## [2,]    7    2    5  
## [3,]    5    3    5
```

```
z <- matrix(y, nrow=3, byrow=T)
```

```
# By changing the "byrow" option, we can fill the matrix by row  
z
```

```
##      [,1] [,2] [,3]  
## [1,]    3    7    5  
## [2,]    1    2    3  
## [3,]    2    5    5
```

```
## Extracting values from matrices:
```

```
z[2,] # Row
```

```
## [1] 1 2 3
```

```
z[,3] # Column
```

```
## [1] 5 3 5
```

```
z[2,3] # Value
```

```
## [1] 3
```

```
## Other ways to make matrices:
```

```
array(data = y, dim = c(3,3))
```

```
##      [,1] [,2] [,3]  
## [1,]    3    1    2  
## [2,]    7    2    5  
## [3,]    5    3    5
```

```
matrix(c(1,2,3,4,5,6), nrow = 2)
```

```
##      [,1] [,2] [,3]  
## [1,]    1    3    5  
## [2,]    2    4    6
```

```
## Create Dataframes:
dat <- as.data.frame(z)
names(dat) <- c("cat", "giraffe", "bowlingball")
dat

##   cat giraffe bowlingball
## 1   3       7           5
## 2   1       2           3
## 3   2       5           5

## R has base functions:
mean(y)

## [1] 3.666667

length(y)

## [1] 9

sd(y)

## [1] 1.936492

var(y)

## [1] 3.75

prod(y) # Takes the product of each element in the vector

## [1] 31500

apply(z, 2, mean) # Very useful in avoiding for loops, also has useful cousins sapply and lapply

## [1] 2.000000 4.666667 4.333333

#####
## Introduction to Statistics with R
#####

## Getting help
# help.start() # Opens html help in web browser (if installed)
# help(help)   # find help on how to use help
# ?help       # Same as above
# help.search("help") # Find all functions that include the word 'help'

## Getting help on particular functions:
help(plot) # You can use the help function, the argument is an R function
?plot # or just a question mark in front of the function you have questions about

## Reading in your data
getwd() # What directory are we in?
```

```
## [1] "/Users/breannechryst/Desktop/Consultant"
```

```
# R will read and write files to the working directory, unless otherwise specified  
setwd("~/Desktop") # You can change your working directory
```

```
## Read in data
```

```
dat <- read.table("http://www.stat.yale.edu/~blc3/IntroR2015/remote_weight.txt",  
                 header=T, sep=" ", row.names=NULL, as.is = TRUE)  
# Read data including headers, data separated by spaces, no row names  
ls() # List all variables stored in memory
```

```
## [1] "dat" "sub.y" "x" "y" "z"
```

```
head(dat) # Shows the first 6 rows of the data
```

```
##   id remote weight gender  
## 1  1      5    151      0  
## 2  2      7    152      0  
## 3  3      3    153      0  
## 4  4      2    165      0  
## 5  5      5    138      0  
## 6  6      0    149      0
```

```
head(dat, 10) # the first 10 rows of the data
```

```
##   id remote weight gender  
## 1  1      5    151      0  
## 2  2      7    152      0  
## 3  3      3    153      0  
## 4  4      2    165      0  
## 5  5      5    138      0  
## 6  6      0    149      0  
## 7  7      5    142      0  
## 8  8      9    139      0  
## 9  9      1    140      0  
## 10 10      8    138      0
```

```
tail(dat) # last 6 rows of the data
```

```
##   id remote weight gender  
## 95  95     30    167      1  
## 96  96     28    154      1  
## 97  97     29    181      1  
## 98  98     34    172      1  
## 99  99     23    159      1  
## 100 100     25    177      1
```

```
## Extracting data from the data frame
```

```
dim(dat) # Find out how many rows and columns in the data set
```

```
## [1] 100  4
```

```
names(dat)           # List all variable names in the dataset
```

```
## [1] "id"      "remote" "weight" "gender"
```

```
str(dat)             # Look at the structure of your data
```

```
## 'data.frame':   100 obs. of  4 variables:
## $ id      : int  1 2 3 4 5 6 7 8 9 10 ...
## $ remote: int  5 7 3 2 5 0 5 9 1 8 ...
## $ weight: int 151 152 153 165 138 149 142 139 140 138 ...
## $ gender: int  0 0 0 0 0 0 0 0 0 0 ...
```

```
dat                 # See the data frame on the screen
```

```
##      id remote weight gender
## 1      1      5    151      0
## 2      2      7    152      0
## 3      3      3    153      0
## 4      4      2    165      0
## 5      5      5    138      0
## 6      6      0    149      0
## 7      7      5    142      0
## 8      8      9    139      0
## 9      9      1    140      0
## 10     10      8    138      0
## 11     11      7    137      0
## 12     12      7    119      0
## 13     13      8    140      0
## 14     14      9    145      0
## 15     15      5    126      0
## 16     16      3    142      0
## 17     17      6    127      0
## 18     18     10    135      0
## 19     19      7    149      0
## 20     20      7    134      0
## 21     21      5    157      0
## 22     22      7    141      0
## 23     23      4    146      0
## 24     24      8    127      0
## 25     25      1    143      0
## 26     26      4    151      0
## 27     27      8    132      0
## 28     28      2    149      0
## 29     29      0    144      0
## 30     30      5    148      0
## 31     31      2    119      0
## 32     32      9    129      0
## 33     33      8    138      0
## 34     34      1    165      0
## 35     35     10    142      0
## 36     36      8    138      0
## 37     37      5    137      0
```

##	38	38	9	154	0
##	39	39	1	133	0
##	40	40	9	139	0
##	41	41	2	146	0
##	42	42	9	152	0
##	43	43	7	149	0
##	44	44	10	128	0
##	45	45	8	131	0
##	46	46	4	136	0
##	47	47	1	148	0
##	48	48	4	134	0
##	49	49	1	137	0
##	50	50	4	137	0
##	51	51	31	159	1
##	52	52	32	167	1
##	53	53	27	175	1
##	54	54	37	189	1
##	55	55	28	179	1
##	56	56	27	145	1
##	57	57	30	161	1
##	58	58	35	181	1
##	59	59	31	177	1
##	60	60	27	178	1
##	61	61	26	170	1
##	62	62	31	158	1
##	63	63	32	179	1
##	64	64	37	192	1
##	65	65	35	183	1
##	66	66	31	184	1
##	67	67	29	169	1
##	68	68	34	182	1
##	69	69	30	190	1
##	70	70	30	195	1
##	71	71	35	182	1
##	72	72	30	174	1
##	73	73	37	184	1
##	74	74	29	188	1
##	75	75	23	150	1
##	76	76	31	178	1
##	77	77	35	186	1
##	78	78	28	170	1
##	79	79	21	149	1
##	80	80	33	203	1
##	81	81	27	172	1
##	82	82	23	160	1
##	83	83	36	179	1
##	84	84	30	195	1
##	85	85	32	160	1
##	86	86	20	147	1
##	87	87	22	144	1
##	88	88	35	178	1
##	89	89	34	168	1
##	90	90	27	157	1
##	91	91	34	184	1

```
## 92  92    18   133    1
## 93  93    31   160    1
## 94  94    29   170    1
## 95  95    30   167    1
## 96  96    28   154    1
## 97  97    29   181    1
## 98  98    34   172    1
## 99  99    23   159    1
## 100 100    25   177    1
```

```
dat[1:10,]      # See the first 10 rows
```

```
##      id remote weight gender
## 1     1      5    151      0
## 2     2      7    152      0
## 3     3      3    153      0
## 4     4      2    165      0
## 5     5      5    138      0
## 6     6      0    149      0
## 7     7      5    142      0
## 8     8      9    139      0
## 9     9      1    140      0
## 10    10     8    138      0
```

```
dat[, "weight"]  # See only the weight column
```

```
##      [1] 151 152 153 165 138 149 142 139 140 138 137 119 140 145 126 142 127
##      [18] 135 149 134 157 141 146 127 143 151 132 149 144 148 119 129 138 165
##      [35] 142 138 137 154 133 139 146 152 149 128 131 136 148 134 137 137 159
##      [52] 167 175 189 179 145 161 181 177 178 170 158 179 192 183 184 169 182
##      [69] 190 195 182 174 184 188 150 178 186 170 149 203 172 160 179 195 160
##      [86] 147 144 178 168 157 184 133 160 170 167 154 181 172 159 177
```

```
dat[, 3]         # Same as above
```

```
##      [1] 151 152 153 165 138 149 142 139 140 138 137 119 140 145 126 142 127
##      [18] 135 149 134 157 141 146 127 143 151 132 149 144 148 119 129 138 165
##      [35] 142 138 137 154 133 139 146 152 149 128 131 136 148 134 137 137 159
##      [52] 167 175 189 179 145 161 181 177 178 170 158 179 192 183 184 169 182
##      [69] 190 195 182 174 184 188 150 178 186 170 149 203 172 160 179 195 160
##      [86] 147 144 178 168 157 184 133 160 170 167 154 181 172 159 177
```

```
dat$weight       # Yet another way
```

```
##      [1] 151 152 153 165 138 149 142 139 140 138 137 119 140 145 126 142 127
##      [18] 135 149 134 157 141 146 127 143 151 132 149 144 148 119 129 138 165
##      [35] 142 138 137 154 133 139 146 152 149 128 131 136 148 134 137 137 159
##      [52] 167 175 189 179 145 161 181 177 178 170 158 179 192 183 184 169 182
##      [69] 190 195 182 174 184 188 150 178 186 170 149 203 172 160 179 195 160
##      [86] 147 144 178 168 157 184 133 160 170 167 154 181 172 159 177
```



```
dat[1:10,"weight"] # See only the first 10 values of the weight col.
```

```
## [1] 151 152 153 165 138 149 142 139 140 138
```

```
dat[,-1] # See all but the first column of data
```

```
##      remote weight gender
## 1         5     151      0
## 2         7     152      0
## 3         3     153      0
## 4         2     165      0
## 5         5     138      0
## 6         0     149      0
## 7         5     142      0
## 8         9     139      0
## 9         1     140      0
## 10        8     138      0
## 11        7     137      0
## 12        7     119      0
## 13        8     140      0
## 14        9     145      0
## 15        5     126      0
## 16        3     142      0
## 17        6     127      0
## 18       10     135      0
## 19        7     149      0
## 20        7     134      0
## 21        5     157      0
## 22        7     141      0
## 23        4     146      0
## 24        8     127      0
## 25        1     143      0
## 26        4     151      0
## 27        8     132      0
## 28        2     149      0
## 29        0     144      0
## 30        5     148      0
## 31        2     119      0
## 32        9     129      0
## 33        8     138      0
## 34        1     165      0
## 35       10     142      0
## 36        8     138      0
## 37        5     137      0
## 38        9     154      0
## 39        1     133      0
## 40        9     139      0
## 41        2     146      0
## 42        9     152      0
## 43        7     149      0
## 44       10     128      0
## 45        8     131      0
```

## 46	4	136	0
## 47	1	148	0
## 48	4	134	0
## 49	1	137	0
## 50	4	137	0
## 51	31	159	1
## 52	32	167	1
## 53	27	175	1
## 54	37	189	1
## 55	28	179	1
## 56	27	145	1
## 57	30	161	1
## 58	35	181	1
## 59	31	177	1
## 60	27	178	1
## 61	26	170	1
## 62	31	158	1
## 63	32	179	1
## 64	37	192	1
## 65	35	183	1
## 66	31	184	1
## 67	29	169	1
## 68	34	182	1
## 69	30	190	1
## 70	30	195	1
## 71	35	182	1
## 72	30	174	1
## 73	37	184	1
## 74	29	188	1
## 75	23	150	1
## 76	31	178	1
## 77	35	186	1
## 78	28	170	1
## 79	21	149	1
## 80	33	203	1
## 81	27	172	1
## 82	23	160	1
## 83	36	179	1
## 84	30	195	1
## 85	32	160	1
## 86	20	147	1
## 87	22	144	1
## 88	35	178	1
## 89	34	168	1
## 90	27	157	1
## 91	34	184	1
## 92	18	133	1
## 93	31	160	1
## 94	29	170	1
## 95	30	167	1
## 96	28	154	1
## 97	29	181	1
## 98	34	172	1
## 99	23	159	1

```
## 100      25      177      1
```

```
dat.o <- dat      # Copy the data frame to a data.frame named dat.o
ls()              # Now we have 5 variables: 'x', 'y', 'z', 'dat' and 'dat.o'
```

```
## [1] "dat"  "dat.o" "sub.y" "x"     "y"     "z"
```

```
## Getting familiar with the data
summary(dat)      # Generate summary statistics of data
```

```
##           id           remote           weight           gender
## Min.      : 1.00    Min.      : 0.00    Min.      :119.0    Min.      :0.0
## 1st Qu.: 25.75    1st Qu.: 5.00    1st Qu.:139.8    1st Qu.:0.0
## Median : 50.50    Median :14.00    Median :152.0    Median :0.5
## Mean     : 50.50    Mean     :17.59    Mean     :156.4    Mean     :0.5
## 3rd Qu.: 75.25    3rd Qu.:30.00    3rd Qu.:174.2    3rd Qu.:1.0
## Max.     :100.00    Max.      :37.00    Max.      :203.0    Max.      :1.0
```

```
apply(dat, 2, sd) # Calculate standard deviations of all variables
```

```
##           id      remote      weight      gender
## 29.0114920 12.8488454 20.0833491 0.5025189
```

```
var(dat)          # Variance on diagonal, covariance off diagonal
```

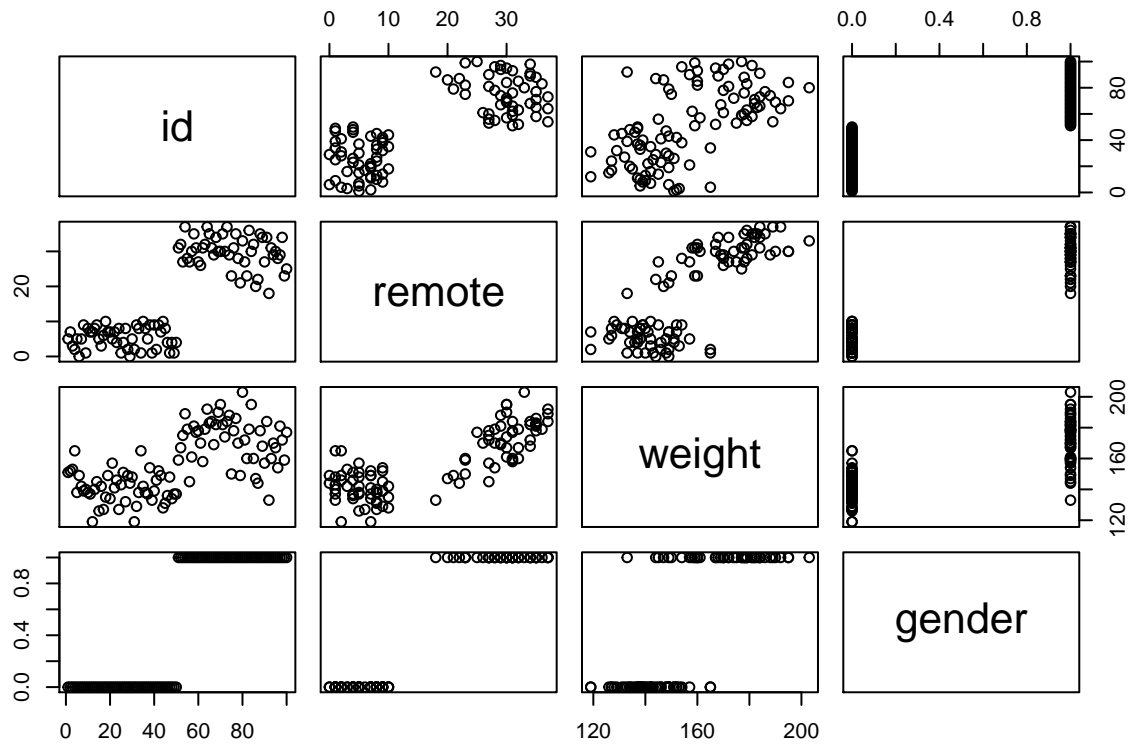
```
##           id      remote      weight      gender
## id      841.66667 299.287879 355.691919 12.6262626
## remote 299.28788 165.092828 209.448990  6.1565657
## weight 355.69192 209.448990 403.340909  7.7929293
## gender 12.62626   6.156566   7.792929  0.2525253
```

```
mean(dat)         # Calculate the mean of all variables
```


```
## Warning in mean.default(dat): argument is not numeric or logical: returning
## NA
```

```
## [1] NA
```


```
pairs(dat)        # A general view of data through scatter plots
```



```
pairs(dat[, -1])
```



```
# See scatterplots for all pairs of variables except the first ('id') in the data frame
plot(remote ~ weight, data = dat) # Scatterplot of 'weight' vs. 'remote'
```



```
# Changing data type
class(dat$gender) # What kind of variable is 'gender'?
```

```
## [1] "integer"
```

```
dat$gender <- factor(dat$gender) # Converts 'gender' from type integer to factor
class(dat$gender) # Verify that gender is now indeed of type factor
```

```
## [1] "factor"
```

```
dat$gender # See all data in column 'gender'; note "Levels: 0 1" at the bottom
```

```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [36] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [71] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## Levels: 0 1
```

```
# Attaching the data frame
attach(dat) # Attach the data frame
remote # Now we can refer directly to the variable without using $
```


```
## [1] 5 7 3 2 5 0 5 9 1 8 7 7 8 9 5 3 6 10 7 7 5 7 4
## [24] 8 1 4 8 2 0 5 2 9 8 1 10 8 5 9 1 9 2 9 7 10 8 4
## [47] 1 4 1 4 31 32 27 37 28 27 30 35 31 27 26 31 32 37 35 31 29 34 30
## [70] 30 35 30 37 29 23 31 35 28 21 33 27 23 36 30 32 20 22 35 34 27 34 18
## [93] 31 29 30 28 29 34 23 25
```

```
# Basic Graphics
```

```
hist(remote)
```

```
# Histogram of 'remote'
```


Histogram of remote




```
hist(weight)
```

```
# Histogram of 'weight'
```


Histogram of weight



```
boxplot(remote,weight)      # Boxplot of 'remote' and 'weight'
```




```
boxplot(remote ~ gender)    # Boxplot of 'remote' conditioned on 'gender'
```



```
boxplot(weight ~ gender, main = "Weight by Gender",
        xlab = "Gender", ylab = "Weight",
        col = c("forestgreen", "tomato")) # Boxplot of 'weight' conditioned on 'gender'
```

Weight by Gender



```
## Inferential statistics
```

```
cor(remote,weight) # Run correlation coefficient
```

```
## [1] 0.8116673
```

```
t.test(remote ~ gender) # Did frequency of remote use differ by gender?
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: remote by gender
```

```
## t = -31.32, df = 84.845, p-value < 2.2e-16
```

```
## alternative hypothesis: true difference in means is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## -25.92777 -22.83223
```

```
## sample estimates:
```

```
## mean in group 0 mean in group 1
```

```
## 5.40 29.78
```

```
rem.t <- t.test(remote ~ gender) # Save results of last analysis
```

```
rem.t # Display analysis
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: remote by gender
```

```
## t = -31.32, df = 84.845, p-value < 2.2e-16
```



```
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -25.92777 -22.83223
## sample estimates:
## mean in group 0 mean in group 1
##          5.40          29.78
```

```
names(rem.t)          # See the names of variables in the object rem.t
```

```
## [1] "statistic" "parameter" "p.value" "conf.int" "estimate"
## [6] "null.value" "alternative" "method" "data.name"
```

```
rem.t$statistic      # See the statistics variable in the object rem.t
```

```
##          t
## -31.31951
```

```
mod1 <- lm(remote ~ gender)      # Linear model, regressing 'remote' on 'gender'
anova(mod1)                      # ANOVA table of the previous model
```

```
## Analysis of Variance Table
##
## Response: remote
##          Df Sum Sq Mean Sq F value    Pr(>F)
## gender      1 14859.6 14859.6  980.91 < 2.2e-16 ***
## Residuals  98  1484.6    15.1
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(lm(remote ~ gender))      # You can combine the two steps in to one line
```

```
## Analysis of Variance Table
##
## Response: remote
##          Df Sum Sq Mean Sq F value    Pr(>F)
## gender      1 14859.6 14859.6  980.91 < 2.2e-16 ***
## Residuals  98  1484.6    15.1
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
mod2 <- lm(remote ~ weight)      # Model 'remote' as a linear function of 'weight'
mod3 <- lm(remote ~ weight + gender) # Model 'remote' as a linear function of 'weight' & 'gender'
mod4 <- lm(remote ~ weight*gender)
# Equivalent to all main effects and interaction:
# lm(remote ~ weight + gender + weight*gender)
summary(mod3)                   # See regression table for model 3 (remote ~ weight + gender)
```

```
##
## Call:
## lm(formula = remote ~ weight + gender)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.2651 -2.6495  0.1842  2.9608  6.1556
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -11.44900     4.02622  -2.844  0.00544 **
## weight       0.11948     0.02832   4.219 5.53e-05 ***
## gender1      20.69286     1.13190  18.282 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.596 on 97 degrees of freedom
## Multiple R-squared:  0.9232, Adjusted R-squared:  0.9217
## F-statistic: 583.4 on 2 and 97 DF,  p-value: < 2.2e-16
```

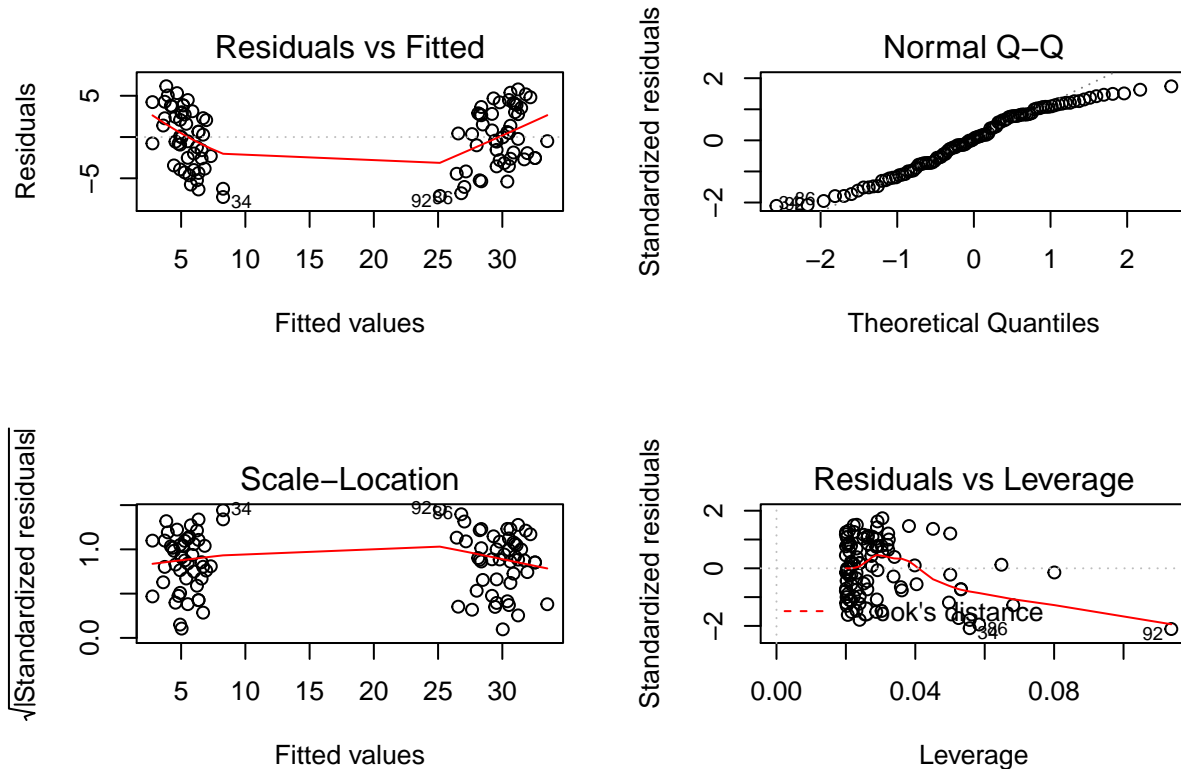
```
summary(mod4)           # See regression table for model 4 (remote ~ weight*gender)
```

```
##
## Call:
## lm(formula = remote ~ weight * gender)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.8530 -2.7707  0.0422  2.5473  5.0332
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   16.72442     6.32943   2.642  0.00962 **
## weight        -0.08030     0.04477  -1.794  0.07601 .
## gender1       -22.96669     8.18338  -2.807  0.00606 **
## weight:gender1  0.28988     0.05393   5.375 5.36e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.169 on 96 degrees of freedom
## Multiple R-squared:  0.941, Adjusted R-squared:  0.9392
## F-statistic: 510.4 on 3 and 96 DF,  p-value: < 2.2e-16
```

```
anova(mod3,mod4)       # Prints ANOVA table comparing model 3 to model 4 (delta F)
```

```
## Analysis of Variance Table
##
## Model 1: remote ~ weight + gender
## Model 2: remote ~ weight * gender
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      97 1254.43
## 2      96  964.23   1    290.2 28.893 5.359e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Regression diagnostics
mod3.1 <- lm(remote ~ weight + gender) # Gives you regression diagnostics
par(mfrow=c(2,2))                     # Set up plotting region for a 2x2 grid
plot(mod3.1)                           # Plot the regression diagnostics (R knows automatically to do this)
```



```
## Saving the graphs as PDF
pdf("prettygraph.pdf") # Turn on the PDF device and open a blank file called "prettygraph.ps"
plot(mod3.1)           # Plot the model
dev.off()              # Turn off the postscript device
```

```
## pdf
## 2
```

```
#####
## Intro to Simulation and Functions with R:
#####
```

```
## Create a vector of 12 random numbers drawn
## from a uniform distribution over the
## interval between 0 and 1:
z <- runif(12) # Generates 12 observations from Unif(0,1)
```

```
z
```

```
## [1] 0.4556848 0.9854712 0.4200237 0.8291504 0.4703348 0.2857762 0.6584719
## [8] 0.2546510 0.9307130 0.5630982 0.6581636 0.0478349
```

```
## We can see which of these is less than 0.5 with the expression "z < 0.5"
z < 0.5
```

```
## [1] TRUE FALSE TRUE FALSE TRUE TRUE FALSE TRUE FALSE FALSE FALSE
## [12] TRUE
```

```
## R identifies "True" with "1" and "False" with "0":
as.numeric(z < 0.5)
```

```
## [1] 1 0 1 0 1 1 0 1 0 0 0 1
```

```
## Now let x be a vector of 100 random draws from a "Normal" distribution:
x <- rnorm(100) # Generates 100 random normal observations, mean 0 sd 1

x
```


```
## [1] -0.60707335 -0.33005169 0.62230989 0.38551562 0.31543422
## [6] 2.47740446 -1.86243466 1.39123147 0.43647982 -0.73471005
## [11] 0.75510995 -1.38845849 0.78355418 1.13443368 1.90415334
## [16] -2.08424796 0.63740099 0.57301829 1.07548641 0.16850846
## [21] 0.37881990 0.95210272 1.56399504 1.94397236 0.62846865
## [26] 0.42433338 -0.01904825 0.12624928 -0.91928159 -0.25292838
## [31] 0.51290226 -0.01157962 -1.17750323 -0.89039806 1.46312738
## [36] -1.42614654 1.17351646 1.91803359 -1.37987164 -0.45621027
## [41] -0.30672590 -0.10949010 0.51073777 0.80178036 0.19224516
## [46] -0.04049865 0.59691799 -0.42127898 -0.42071509 -0.30627443
## [51] 1.75439433 0.55773478 0.75073092 -0.35571323 -0.55421789
## [56] 2.06204030 -1.02788681 0.32160627 -1.78920299 0.28950207
## [61] -0.17114089 1.38651570 0.52806083 0.77039366 0.51529084
## [66] 1.29652106 -0.68216217 -1.08444008 -0.93716872 -0.77377716
## [71] 0.69587054 0.81513938 -1.07469060 0.12932937 0.11704900
## [76] 0.75277881 -0.40551496 -0.37512544 -0.46837951 0.21524363
## [81] 0.27444701 -0.08730682 -1.48489614 0.85683601 -1.33118827
## [86] 0.83443599 1.36535987 0.01713701 0.81882825 -0.03155508
## [91] 0.88367911 0.66498246 -1.75013434 0.55482368 0.02422135
## [96] 0.99295549 1.74252547 0.47925919 1.00161752 0.32721646
```

```
par(mfrow=c(1,1)) #Resets the graphics to one plot per page
hist(x, main= "100 Obs. Standard Normal Distribution",
     breaks = 10, col = "gold") # Create a histogram of the vector x
```

100 Obs. Standard Normal Distribution




```
y <- rnorm(100) # save 100 random sample from a standard normal distribution to y  
plot(x,y)
```



```
plot(x,y,main = "Noise", col = "steelblue", pch = 18)  
abline(h=0)  
legend("topleft", title = "I'm Normal!", "Obs", pch= 18,
```


```
col = "steelblue")
```

Noise



```
plot(x,y,type="l", main = "Nonsense", col = "coral")
```

Nonsense



```
##plot
```

```
# Sampling uniformly at random, with replacement
```

```
v <- sample(1:10,100,replace=T) # Samples from 1 to 10, 100 times
```

```
v
```

```
## [1] 10 7 1 3 1 10 6 6 9 8 2 5 6 1 3 8 3 8 4 6 4 9 7
## [24] 7 8 8 9 7 10 4 5 7 6 8 1 5 2 7 4 2 1 10 10 8 10 5
## [47] 3 9 3 2 8 5 9 1 3 6 10 1 10 5 6 7 5 8 1 6 1 1 2
## [70] 2 7 9 1 1 3 1 4 7 1 4 2 6 1 7 1 4 7 8 1 4 6 2
## [93] 10 4 2 2 3 6 3 10
```

```
table(v)
```

```
## v
```

```
## 1 2 3 4 5 6 7 8 9 10
```

```
## 17 10 9 9 7 11 11 10 6 10
```

```
## Functions:
```

```
## Numerical calculations for birthday problem:
```

```
k <- 40
```

```
top <- seq(365,length=k,by=-1) # Creates a vector of 365 to 365-k
```

```
bottom <- rep(365,k) # Creates a vector filled with 365 repeated k times
```

```
top
```

```
## [1] 365 364 363 362 361 360 359 358 357 356 355 354 353 352 351 350 349
```

```
## [18] 348 347 346 345 344 343 342 341 340 339 338 337 336 335 334 333 332
```

```
## [35] 331 330 329 328 327 326
```

```
bottom
```

```
## [1] 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365
```

```
## [18] 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365
```

```
## [35] 365 365 365 365 365 365
```

```
top/bottom
```

```
## [1] 1.0000000 0.9972603 0.9945205 0.9917808 0.9890411 0.9863014 0.9835616
```

```
## [8] 0.9808219 0.9780822 0.9753425 0.9726027 0.9698630 0.9671233 0.9643836
```

```
## [15] 0.9616438 0.9589041 0.9561644 0.9534247 0.9506849 0.9479452 0.9452055
```

```
## [22] 0.9424658 0.9397260 0.9369863 0.9342466 0.9315068 0.9287671 0.9260274
```

```
## [29] 0.9232877 0.9205479 0.9178082 0.9150685 0.9123288 0.9095890 0.9068493
```

```
## [36] 0.9041096 0.9013699 0.8986301 0.8958904 0.8931507
```

```
prod(top/bottom) # This is the prob of NO birthday match
```

```
## [1] 0.1087682
```

```
1 - prod(top/bottom) # This is the prob of having a birthday match
```

```
## [1] 0.8912318
```

```
## Let's make a function out of what we just did:
```

```
bday <- function(k){ # k is the variable  
  top <- seq(365,length=k,by=-1)  
  bottom <- rep(365,k)  
  return(1-prod(top/bottom))  
}
```

```
bday(40)
```

```
## [1] 0.8912318
```

```
#####  
## Intro to for loops in R:  
#####
```

```
s <- 0  
for(i in 1:100){  
  s <- s+i  
}  
s
```

```
## [1] 5050
```

```
## Sometimes you can do the same thing without a loop:  
sum(1:100)
```

```
## [1] 5050
```

```
## You can have more commands in the body of the loop:
```

```
s <- 0  
for(i in 1:10){  
  s <- s+i  
  cat("When i = ", i, ", s = ", s, "\n", sep="") # "cat" prints things  
}
```

```
## When i = 1, s = 1  
## When i = 2, s = 3  
## When i = 3, s = 6  
## When i = 4, s = 10  
## When i = 5, s = 15  
## When i = 6, s = 21
```



```
## When i = 7, s = 28
## When i = 8, s = 36
## When i = 9, s = 45
## When i = 10, s = 55
```

```
s <- 0
for(i in 1:10){
  s <- s+i
  cat("When i = ", i, ", s = ", s, "\n", sep="")
  remainder2 <- (i %% 2)
  twos <- i/2
  if(remainder2 == 0){
    cat("I'm getting", rep("really", twos), "tired!\n")
  }
}
```

```
## When i = 1, s = 1
## When i = 2, s = 3
## I'm getting really tired!
## When i = 3, s = 6
## When i = 4, s = 10
## I'm getting really really tired!
## When i = 5, s = 15
## When i = 6, s = 21
## I'm getting really really really tired!
## When i = 7, s = 28
## When i = 8, s = 36
## I'm getting really really really really tired!
## When i = 9, s = 45
## When i = 10, s = 55
## I'm getting really really really really really tired!
```