

Intro__to__R__2016__Code.R

breannechryst

Tue Apr 5 17:21:57 2016

```
## Introduction to R
## Breanne Chryst
## CSSSI StatLab
## April 6, 2016
```

```
#####
## R Basics
#####
```

```
## R can be used as a calculator, it works as expected:
2+3
```

```
## [1] 5
```

```
exp(2)
```

```
## [1] 7.389056
```

```
5^(2)
```

```
## [1] 25
```

```
3*2+4
```

```
## [1] 10
```

```
(2*4+1)/2
```

```
## [1] 4.5
```

```
## The print command:
print("Hello World")
```

```
## [1] "Hello World"
```

```
print(5*2-11)
```

```
## [1] -1
```

```
## Assigning a variable
```

```
x <- 5 # 5 has now been assigned to the variable x
```

```
# the "<-" assigns the value on the right to the name on the left. Made by "alt -"
```

```
x # Evaluation
```

```
## [1] 5
```

```
x^2
```

```
## [1] 25
```

```
## Creating a vector:
```

```
y <- c(3,7,5,1,2,3,2,5,5) # "c()" concatenates, creating a vector
```

```
## Extracting values of a vector:
```

```
y[2]
```

```
## [1] 7
```

```
3:5 # the whole numbers from 3 to 5
```

```
## [1] 3 4 5
```

```
y[3:5] # extracts the 3rd to 5th elements in the vector y
```

```
## [1] 5 1 2
```

```
sub.y <- y[3:5]
```

```
sub.y
```

```
## [1] 5 1 2
```

```
## Value Comparison
```

```
2==2 # Equality
```

```
## [1] TRUE
```

```
2!=2 # Inequality
```

```
## [1] FALSE
```

```
x <= y # less than or equal: "<", ">", and ">=" also work
```

```
## [1] FALSE TRUE TRUE FALSE FALSE FALSE FALSE TRUE TRUE
```

```
## Other ways to make vectors:
```

```
array(data = 0, dim = 3)
```

```
## [1] 0 0 0
```

```
seq(from = 1, to = 4, by = 1)
```

```
## [1] 1 2 3 4
```

```
seq(1,4,0.5)
```

```
## [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0
```

```
rep(x = "cat", times = 3)
```

```
## [1] "cat" "cat" "cat"
```

```
rep(c("cat", 4, x^(2)), each = 2)
```

```
## [1] "cat" "cat" "4" "4" "25" "25"
```

```
rep(c("cat", 4, x^(2)), times = 2)
```

```
## [1] "cat" "4" "25" "cat" "4" "25"
```

```
## "matrix()" creates a matrix from the values entered:
```

```
z <- matrix(y, nrow=3) # This is filled by column  
z
```

```
##      [,1] [,2] [,3]  
## [1,]    3    1    2  
## [2,]    7    2    5  
## [3,]    5    3    5
```

```
z <- matrix(y, nrow=3, byrow=T)  
# By changing the "byrow" option, we can fill the matrix by row  
z
```

```
##      [,1] [,2] [,3]  
## [1,]    3    7    5  
## [2,]    1    2    3  
## [3,]    2    5    5
```

```
## Extracting values from matrices:
```

```
z[2,] # Row
```

```
## [1] 1 2 3
```

```
z[,3] # Column
```

```
## [1] 5 3 5
```

```
z[2,3] # Value
```

```
## [1] 3
```

```
## Other ways to make matrices:
```

```
array(data = y, dim = c(3,3))
```

```
##      [,1] [,2] [,3]
## [1,]    3    1    2
## [2,]    7    2    5
## [3,]    5    3    5
```

```
matrix(c(1,2,3,4,5,6), nrow = 2)
```

```
##      [,1] [,2] [,3]
## [1,]    1    3    5
## [2,]    2    4    6
```

```
## Create Dataframes:
```

```
dat <- as.data.frame(z)
```

```
names(dat) <- c("cat", "giraffe", "bowlingball")
```

```
dat
```

```
##   cat giraffe bowlingball
## 1   3       7           5
## 2   1       2           3
## 3   2       5           5
```

```
dat2 <- data.frame( ID=1:4,
                    FirstName=c("Batman", "Robin", "Ivy", "Joker"),
                    Female=c(F,F,T,F),
                    Age=c(22,33,44,55) )
```

```
dat2$FirstName # Access the second column of dfr1.
```

```
## [1] Batman Robin Ivy   Joker
## Levels: Batman Ivy Joker Robin
```

```
## R has base functions:
```

```
mean(y)
```

```
## [1] 3.666667
```

```
length(y)
```

```
## [1] 9
```

```
sd(y)
```

```
## [1] 1.936492
```

```
var(y)
```

```
## [1] 3.75
```

```
prod(y) # Takes the product of each element in the vector
```

```
## [1] 31500
```

```
cor(z)
```

```
##           [,1]      [,2]      [,3]
## [1,] 1.0000000 0.9933993 0.8660254
## [2,] 0.9933993 1.0000000 0.9176629
## [3,] 0.8660254 0.9176629 1.0000000
```

```
apply(z, 2, mean) # Very useful in avoiding for loops, also has useful cousins sapply and lapply
```

```
## [1] 2.000000 4.666667 4.333333
```

```
#####
## Introduction to Statistics with R
#####
```

```
## Getting help
```

```
# help.start() # Opens html help in web browser (if installed)
```

```
# help(help)   # find help on how to use help
```

```
# ?help        # Same as above
```

```
# help.search("help") # Find all functions that include the word 'help'
```

```
## Getting help on particular functions:
```

```
help(plot) # You can use the help function, the argument is an R function
```

```
?plot # or just a question mark in front of the function you have questions about
```

```
## Reading in your data
```

```
getwd()          # What directory are we in?
```

```
## [1] "/Users/breannechryst/IntroductionToR2016"
```

```
# R will read and write files to the working directory, unless otherwise specified
setwd("~/IntroductionToR2016") # You can change your working directory
```

```
## Read in data from the web
```

```
dat <- read.table("http://www.stat.yale.edu/~blc3/IntroR2016/remote_weight.txt",
                  header=T, sep="", row.names=NULL, as.is = TRUE)
```

```
## Read in Data from local folder
dat <- read.table("~/IntroductionToR2016/remote_weight.txt",
                  header=T, sep=" ", as.is = TRUE)

# Read data including headers, data separated by spaces, no row names
ls()                                # List all variables stored in memory

## [1] "dat"      "dat2"      "sub.y"     "x"         "y"         "z"

head(dat)                          # Shows the first 6 rows of the data

##      id remote weight gender group
## 1  1      5     151      0      1
## 2  2      7     152      0      2
## 3  3      3     153      0      1
## 4  4      2     165      0      2
## 5  5      5     138      0      1
## 6  6      0     149      0      2

head(dat, 10)                      # the first 10 rows of the data

##      id remote weight gender group
## 1  1      5     151      0      1
## 2  2      7     152      0      2
## 3  3      3     153      0      1
## 4  4      2     165      0      2
## 5  5      5     138      0      1
## 6  6      0     149      0      2
## 7  7      5     142      0      1
## 8  8      9     139      0      2
## 9  9      1     140      0      2
## 10 10      8     138      0      1

tail(dat)                          # last 6 rows of the data

##      id remote weight gender group
## 95  95      30     167      1      1
## 96  96      28     154      1      1
## 97  97      29     181      1      1
## 98  98      34     172      1      1
## 99  99      23     159      1      2
## 100 100      25     177      1      1

## Extracting data from the data frame
dim(dat)                          # Find out how many rows and columns in the data set

## [1] 100  5
```

```
names(dat)           # List all variable names in the dataset
```

```
## [1] "id"      "remote" "weight" "gender" "group"
```

```
str(dat)             # Look at the structure of your data
```

```
## 'data.frame':   100 obs. of  5 variables:
## $ id      : int  1 2 3 4 5 6 7 8 9 10 ...
## $ remote: int  5 7 3 2 5 0 5 9 1 8 ...
## $ weight: int 151 152 153 165 138 149 142 139 140 138 ...
## $ gender: int  0 0 0 0 0 0 0 0 0 0 ...
## $ group  : int  1 2 1 2 1 2 1 2 2 1 ...
```

```
dat                  # See the data frame on the screen
```

```
##      id remote weight gender group
## 1      1      5    151      0      1
## 2      2      7    152      0      2
## 3      3      3    153      0      1
## 4      4      2    165      0      2
## 5      5      5    138      0      1
## 6      6      0    149      0      2
## 7      7      5    142      0      1
## 8      8      9    139      0      2
## 9      9      1    140      0      2
## 10    10      8    138      0      1
## 11    11      7    137      0      1
## 12    12      7    119      0      1
## 13    13      8    140      0      2
## 14    14      9    145      0      1
## 15    15      5    126      0      2
## 16    16      3    142      0      2
## 17    17      6    127      0      2
## 18    18     10    135      0      1
## 19    19      7    149      0      2
## 20    20      7    134      0      2
## 21    21      5    157      0      2
## 22    22      7    141      0      2
## 23    23      4    146      0      2
## 24    24      8    127      0      1
## 25    25      1    143      0      2
## 26    26      4    151      0      2
## 27    27      8    132      0      1
## 28    28      2    149      0      1
## 29    29      0    144      0      2
## 30    30      5    148      0      1
## 31    31      2    119      0      2
## 32    32      9    129      0      2
## 33    33      8    138      0      2
## 34    34      1    165      0      1
## 35    35     10    142      0      1
## 36    36      8    138      0      1
```

## 37	37	5	137	0	2
## 38	38	9	154	0	1
## 39	39	1	133	0	1
## 40	40	9	139	0	2
## 41	41	2	146	0	2
## 42	42	9	152	0	1
## 43	43	7	149	0	2
## 44	44	10	128	0	2
## 45	45	8	131	0	1
## 46	46	4	136	0	2
## 47	47	1	148	0	2
## 48	48	4	134	0	2
## 49	49	1	137	0	1
## 50	50	4	137	0	1
## 51	51	31	159	1	2
## 52	52	32	167	1	2
## 53	53	27	175	1	1
## 54	54	37	189	1	1
## 55	55	28	179	1	1
## 56	56	27	145	1	1
## 57	57	30	161	1	2
## 58	58	35	181	1	2
## 59	59	31	177	1	1
## 60	60	27	178	1	2
## 61	61	26	170	1	2
## 62	62	31	158	1	1
## 63	63	32	179	1	1
## 64	64	37	192	1	1
## 65	65	35	183	1	2
## 66	66	31	184	1	1
## 67	67	29	169	1	1
## 68	68	34	182	1	2
## 69	69	30	190	1	1
## 70	70	30	195	1	2
## 71	71	35	182	1	2
## 72	72	30	174	1	1
## 73	73	37	184	1	2
## 74	74	29	188	1	2
## 75	75	23	150	1	1
## 76	76	31	178	1	1
## 77	77	35	186	1	2
## 78	78	28	170	1	1
## 79	79	21	149	1	1
## 80	80	33	203	1	2
## 81	81	27	172	1	2
## 82	82	23	160	1	2
## 83	83	36	179	1	2
## 84	84	30	195	1	1
## 85	85	32	160	1	1
## 86	86	20	147	1	1
## 87	87	22	144	1	2
## 88	88	35	178	1	1
## 89	89	34	168	1	1
## 90	90	27	157	1	1


```
## 91  91    34   184    1    1
## 92  92    18   133    1    2
## 93  93    31   160    1    1
## 94  94    29   170    1    2
## 95  95    30   167    1    1
## 96  96    28   154    1    1
## 97  97    29   181    1    1
## 98  98    34   172    1    1
## 99  99    23   159    1    2
## 100 100    25   177    1    1
```

```
dat[1:10,]      # See the first 10 rows
```

```
##      id remote weight gender group
## 1     1      5    151      0      1
## 2     2      7    152      0      2
## 3     3      3    153      0      1
## 4     4      2    165      0      2
## 5     5      5    138      0      1
## 6     6      0    149      0      2
## 7     7      5    142      0      1
## 8     8      9    139      0      2
## 9     9      1    140      0      2
## 10    10     8    138      0      1
```

```
dat[, "weight"] # See only the weight column
```

```
##      [1] 151 152 153 165 138 149 142 139 140 138 137 119 140 145 126 142 127
##      [18] 135 149 134 157 141 146 127 143 151 132 149 144 148 119 129 138 165
##      [35] 142 138 137 154 133 139 146 152 149 128 131 136 148 134 137 137 159
##      [52] 167 175 189 179 145 161 181 177 178 170 158 179 192 183 184 169 182
##      [69] 190 195 182 174 184 188 150 178 186 170 149 203 172 160 179 195 160
##      [86] 147 144 178 168 157 184 133 160 170 167 154 181 172 159 177
```

```
dat[, 3]        # Same as above
```

```
##      [1] 151 152 153 165 138 149 142 139 140 138 137 119 140 145 126 142 127
##      [18] 135 149 134 157 141 146 127 143 151 132 149 144 148 119 129 138 165
##      [35] 142 138 137 154 133 139 146 152 149 128 131 136 148 134 137 137 159
##      [52] 167 175 189 179 145 161 181 177 178 170 158 179 192 183 184 169 182
##      [69] 190 195 182 174 184 188 150 178 186 170 149 203 172 160 179 195 160
##      [86] 147 144 178 168 157 184 133 160 170 167 154 181 172 159 177
```

```
dat$weight      # Yet another way
```

```
##      [1] 151 152 153 165 138 149 142 139 140 138 137 119 140 145 126 142 127
##      [18] 135 149 134 157 141 146 127 143 151 132 149 144 148 119 129 138 165
##      [35] 142 138 137 154 133 139 146 152 149 128 131 136 148 134 137 137 159
##      [52] 167 175 189 179 145 161 181 177 178 170 158 179 192 183 184 169 182
##      [69] 190 195 182 174 184 188 150 178 186 170 149 203 172 160 179 195 160
##      [86] 147 144 178 168 157 184 133 160 170 167 154 181 172 159 177
```

```
dat[1:10,"weight"] # See only the first 10 values of the weight col.
```

```
## [1] 151 152 153 165 138 149 142 139 140 138
```

```
dat[,-1] # See all but the first column of data
```

```
##      remote weight gender group
## 1         5     151      0      1
## 2         7     152      0      2
## 3         3     153      0      1
## 4         2     165      0      2
## 5         5     138      0      1
## 6         0     149      0      2
## 7         5     142      0      1
## 8         9     139      0      2
## 9         1     140      0      2
## 10        8     138      0      1
## 11        7     137      0      1
## 12        7     119      0      1
## 13        8     140      0      2
## 14        9     145      0      1
## 15        5     126      0      2
## 16        3     142      0      2
## 17        6     127      0      2
## 18       10     135      0      1
## 19        7     149      0      2
## 20        7     134      0      2
## 21        5     157      0      2
## 22        7     141      0      2
## 23        4     146      0      2
## 24        8     127      0      1
## 25        1     143      0      2
## 26        4     151      0      2
## 27        8     132      0      1
## 28        2     149      0      1
## 29        0     144      0      2
## 30        5     148      0      1
## 31        2     119      0      2
## 32        9     129      0      2
## 33        8     138      0      2
## 34        1     165      0      1
## 35       10     142      0      1
## 36        8     138      0      1
## 37        5     137      0      2
## 38        9     154      0      1
## 39        1     133      0      1
## 40        9     139      0      2
## 41        2     146      0      2
## 42        9     152      0      1
## 43        7     149      0      2
## 44       10     128      0      2
## 45        8     131      0      1
```

## 46	4	136	0	2
## 47	1	148	0	2
## 48	4	134	0	2
## 49	1	137	0	1
## 50	4	137	0	1
## 51	31	159	1	2
## 52	32	167	1	2
## 53	27	175	1	1
## 54	37	189	1	1
## 55	28	179	1	1
## 56	27	145	1	1
## 57	30	161	1	2
## 58	35	181	1	2
## 59	31	177	1	1
## 60	27	178	1	2
## 61	26	170	1	2
## 62	31	158	1	1
## 63	32	179	1	1
## 64	37	192	1	1
## 65	35	183	1	2
## 66	31	184	1	1
## 67	29	169	1	1
## 68	34	182	1	2
## 69	30	190	1	1
## 70	30	195	1	2
## 71	35	182	1	2
## 72	30	174	1	1
## 73	37	184	1	2
## 74	29	188	1	2
## 75	23	150	1	1
## 76	31	178	1	1
## 77	35	186	1	2
## 78	28	170	1	1
## 79	21	149	1	1
## 80	33	203	1	2
## 81	27	172	1	2
## 82	23	160	1	2
## 83	36	179	1	2
## 84	30	195	1	1
## 85	32	160	1	1
## 86	20	147	1	1
## 87	22	144	1	2
## 88	35	178	1	1
## 89	34	168	1	1
## 90	27	157	1	1
## 91	34	184	1	1
## 92	18	133	1	2
## 93	31	160	1	1
## 94	29	170	1	2
## 95	30	167	1	1
## 96	28	154	1	1
## 97	29	181	1	1
## 98	34	172	1	1
## 99	23	159	1	2

```
## 100      25      177      1      1

dat.o <- dat      # Copy the data frame to a data.frame named dat.o
ls()              # Now we have 5 variables: 'x', 'y', 'z', 'dat' and 'dat.o'

## [1] "dat"      "dat.o" "dat2"  "sub.y" "x"      "y"      "z"

table(dat$group)

##
##  1  2
## 51 49

## Getting familiar with the data
summary(dat)      # Generate summary statistics of data

##           id           remote           weight           gender
## Min.      : 1.00   Min.      : 0.00   Min.      :119.0   Min.      :0.0
## 1st Qu.: 25.75   1st Qu.:  5.00   1st Qu.:139.8   1st Qu.:0.0
## Median : 50.50   Median :14.00   Median :152.0   Median :0.5
## Mean     : 50.50   Mean      :17.59   Mean     :156.4   Mean     :0.5
## 3rd Qu.: 75.25   3rd Qu.:30.00   3rd Qu.:174.2   3rd Qu.:1.0
## Max.     :100.00   Max.      :37.00   Max.      :203.0   Max.      :1.0
##           group
## Min.      :1.00
## 1st Qu.:1.00
## Median :1.00
## Mean     :1.49
## 3rd Qu.:2.00
## Max.     :2.00

apply(dat, 2, sd) # Calculate standard deviations of all variables

##           id      remote      weight      gender      group
## 29.0114920 12.8488454 20.0833491  0.5025189  0.5024184

var(dat)          # Variance on diagonal, covariance off diagonal

##           id      remote      weight      gender      group
## id      841.666667 299.287878 355.691919 12.62626263 -2.24747475
## remote 299.287879 165.0928283 209.4489899  6.15656566 -0.98898990
## weight 355.691919 209.4489899 403.3409091  7.79292929 -0.82878788
## gender 12.626263  6.1565657  7.7929293  0.25252525 -0.03535354
## group  -2.247475 -0.9889899 -0.8287879 -0.03535354  0.25242424

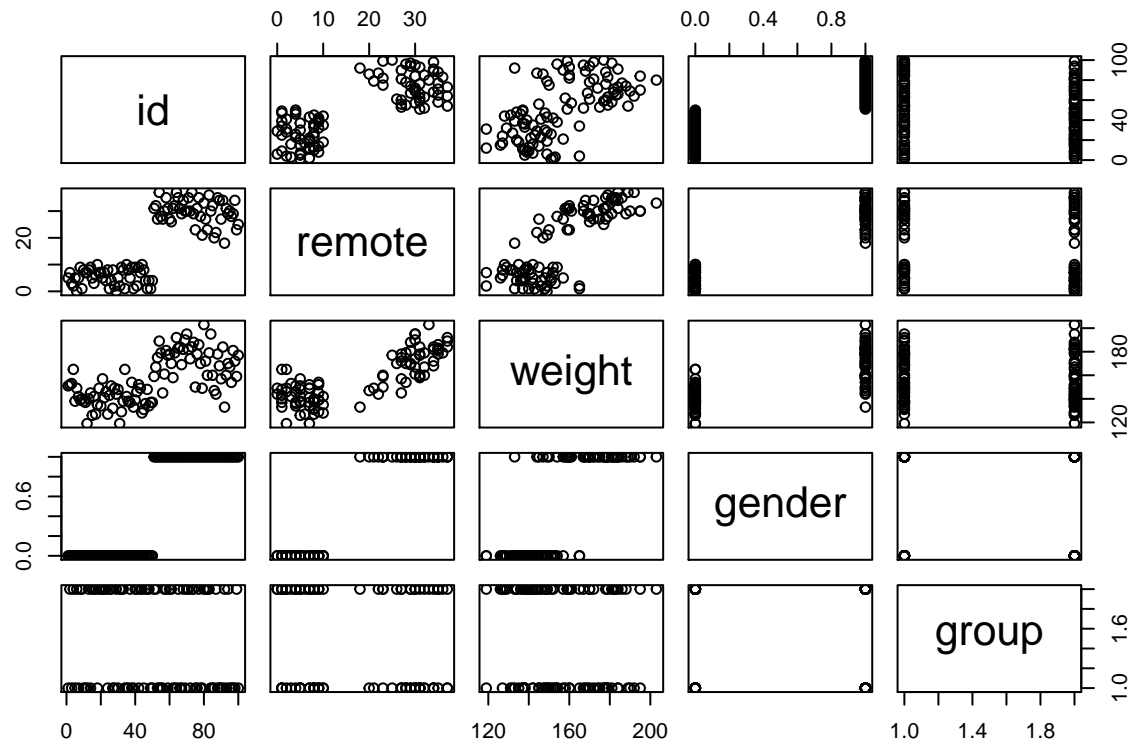
mean(dat)         # Calculate the mean of all variables

## Warning in mean.default(dat): argument is not numeric or logical: returning
## NA

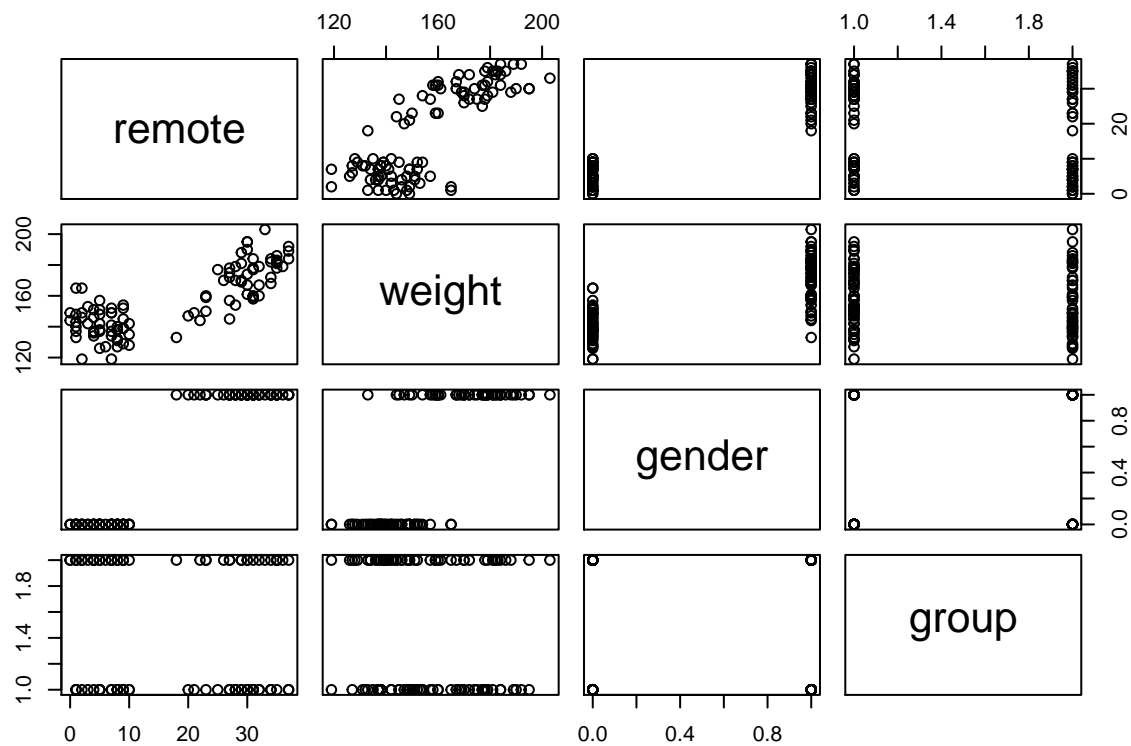
## [1] NA
```

```
pairs(dat)
```


A general view of data through scatter plots



```
pairs(dat[, -1])
```



```
# See scatterplots for all pairs of variables except the first ('id') in the data frame
plot(remote ~ weight, data = dat) # Scatterplot of 'weight' vs. 'remote'
```



```
## Changing data type
class(dat$gender) # What kind of variable is 'gender'?
```

```
## [1] "integer"
```

```
dat$gender <- factor(dat$gender) # Converts 'gender' from type integer to factor
class(dat$gender) # Verify that gender is now indeed of type factor
```

```
## [1] "factor"
```

```
dat$gender # See all data in column 'gender'; note "Levels: 0 1" at the bottom
```

```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [36] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [71] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## Levels: 0 1
```


```
## Attaching the data frame
attach(dat) # Attach the data frame
remote # Now we can refer directly to the variable without using $
```

```
## [1] 5 7 3 2 5 0 5 9 1 8 7 7 8 9 5 3 6 10 7 7 5 7 4
## [24] 8 1 4 8 2 0 5 2 9 8 1 10 8 5 9 1 9 2 9 7 10 8 4
## [47] 1 4 1 4 31 32 27 37 28 27 30 35 31 27 26 31 32 37 35 31 29 34 30
## [70] 30 35 30 37 29 23 31 35 28 21 33 27 23 36 30 32 20 22 35 34 27 34 18
## [93] 31 29 30 28 29 34 23 25
```

```
## Basic Graphics  
hist(remote)
```

```
# Histogram of 'remote'
```


Histogram of remote




```
hist(weight)
```

```
# Histogram of 'weight'
```


Histogram of weight



```
boxplot(remote,weight)      # Boxplot of 'remote' and 'weight'
```




```
boxplot(remote ~ gender)    # Boxplot of 'remote' conditioned on 'gender'
```



```
boxplot(weight ~ gender, main = "Weight by Gender",  
        xlab = "Gender", ylab = "Weight",  
        col = c("forestgreen", "tomato"))      # Boxplot of 'weight' conditioned on 'gender'
```


Weight by Gender



```
## Inferential statistics
```

```
cor(remote,weight) # Run correlation coefficient
```

```
## [1] 0.8116673
```

```
t.test(remote ~ gender) # Did frequency of remote use differ by gender?
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: remote by gender
```

```
## t = -31.32, df = 84.845, p-value < 2.2e-16
```

```
## alternative hypothesis: true difference in means is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## -25.92777 -22.83223
```

```
## sample estimates:
```

```
## mean in group 0 mean in group 1
```

```
## 5.40 29.78
```

```
rem.t <- t.test(remote ~ gender) # Save results of last analysis
```

```
rem.t # Display analysis
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: remote by gender
```

```
## t = -31.32, df = 84.845, p-value < 2.2e-16
```

```
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -25.92777 -22.83223
## sample estimates:
## mean in group 0 mean in group 1
##          5.40          29.78
```

```
names(rem.t)          # See the names of variables in the object rem.t
```

```
## [1] "statistic" "parameter" "p.value" "conf.int" "estimate"
## [6] "null.value" "alternative" "method" "data.name"
```

```
rem.t$statistic      # See the statistics variable in the object rem.t
```

```
##          t
## -31.31951
```

```
mod1 <- lm(remote ~ gender)      # Linear model, regressing 'remote' on 'gender'
anova(mod1)                      # ANOVA table of the previous model
```

```
## Analysis of Variance Table
##
## Response: remote
##          Df Sum Sq Mean Sq F value    Pr(>F)
## gender      1 14859.6 14859.6  980.91 < 2.2e-16 ***
## Residuals  98  1484.6    15.1
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(lm(remote ~ gender))      # You can combine the two steps in to one line
```

```
## Analysis of Variance Table
##
## Response: remote
##          Df Sum Sq Mean Sq F value    Pr(>F)
## gender      1 14859.6 14859.6  980.91 < 2.2e-16 ***
## Residuals  98  1484.6    15.1
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
mod2 <- lm(remote ~ weight)      # Model 'remote' as a linear function of 'weight'
mod3 <- lm(remote ~ weight + gender) # Model 'remote' as a linear function of 'weight' & 'gender'
mod4 <- lm(remote ~ weight*gender)
# Equivalent to all main effects and interaction:
# lm(remote ~ weight + gender + weight*gender)
summary(mod3)                   # See regression table for model 3 (remote ~ weight + gender)
```

```
##
## Call:
## lm(formula = remote ~ weight + gender)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.2651 -2.6495  0.1842  2.9608  6.1556
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -11.44900     4.02622  -2.844  0.00544 **
## weight       0.11948     0.02832   4.219 5.53e-05 ***
## gender1      20.69286     1.13190  18.282 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.596 on 97 degrees of freedom
## Multiple R-squared:  0.9232, Adjusted R-squared:  0.9217
## F-statistic: 583.4 on 2 and 97 DF,  p-value: < 2.2e-16
```

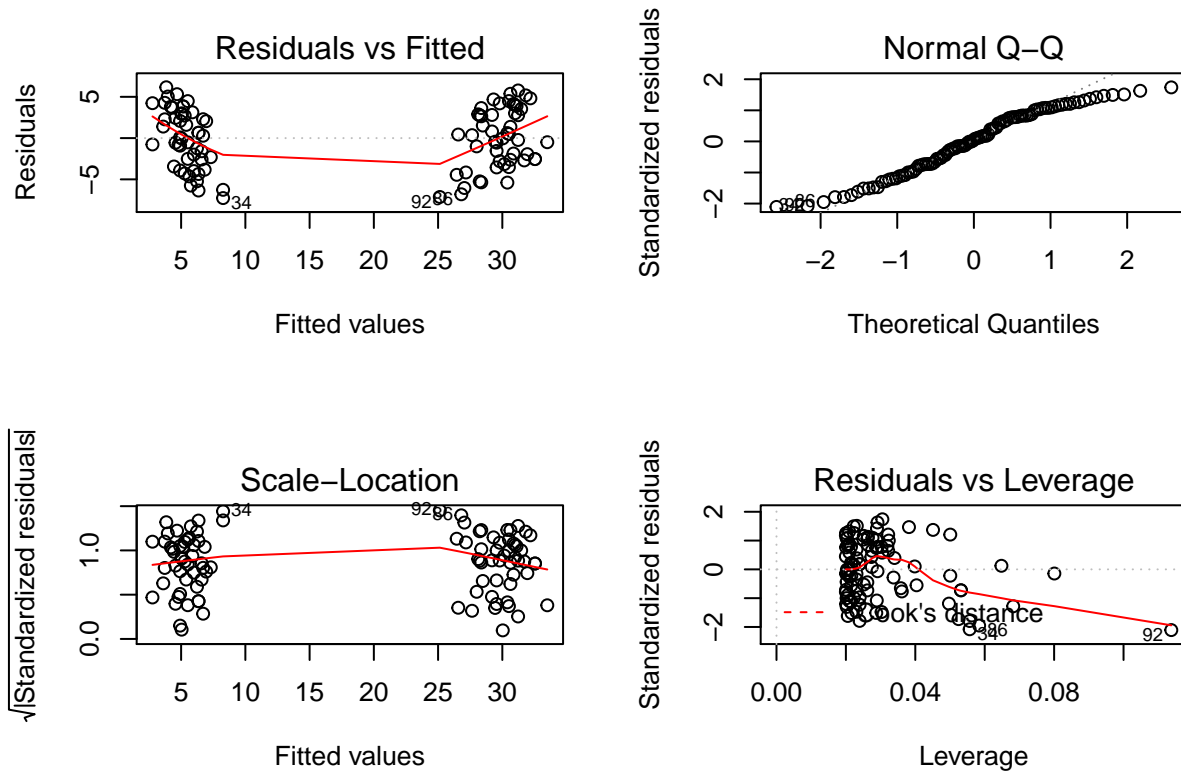
```
summary(mod4)           # See regression table for model 4 (remote ~ weight*gender)
```

```
##
## Call:
## lm(formula = remote ~ weight * gender)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.8530 -2.7707  0.0422  2.5473  5.0332
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   16.72442     6.32943   2.642  0.00962 **
## weight        -0.08030     0.04477  -1.794  0.07601 .
## gender1       -22.96669     8.18338  -2.807  0.00606 **
## weight:gender1  0.28988     0.05393   5.375 5.36e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.169 on 96 degrees of freedom
## Multiple R-squared:  0.941, Adjusted R-squared:  0.9392
## F-statistic: 510.4 on 3 and 96 DF,  p-value: < 2.2e-16
```

```
anova(mod3,mod4)       # Prints ANOVA table comparing model 3 to model 4 (delta F)
```

```
## Analysis of Variance Table
##
## Model 1: remote ~ weight + gender
## Model 2: remote ~ weight * gender
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      97 1254.43
## 2      96  964.23   1    290.2 28.893 5.359e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
## Regression diagnostics
mod3.1 <- lm(remote ~ weight + gender)
par(mfrow=c(2,2))      # Set up plotting region for a 2x2 grid
plot(mod3.1)           # Plot the regression diagnostics (R knows automatically to do this)
```



```
## Saving the graphs as PDF
pdf("prettygraph.pdf") # Turn on the PDF device and open a blank file called "prettygraph.ps"
plot(mod3.1)           # Plot the model
dev.off()              # Turn off the postscript device
```

```
## pdf
## 2
```

```
## Chi-squared Test
chisq.test(gender, group)
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data: gender and group
## X-squared = 1.4406, df = 1, p-value = 0.23
```

```
#####
## Intro to Simulation and Functions with R:
#####

## Create a vector of 12 random numbers drawn
```

```
## from a uniform distribution over the
## interval between 0 and 1:
z <- runif(12) # Generates 12 observations from Unif(0,1)

z

## [1] 0.32920506 0.01898433 0.93448939 0.17891582 0.72237156 0.25391211
## [7] 0.74287966 0.31849095 0.48836747 0.41196728 0.61525964 0.96487975

## We can see which of these is less than 0.5 with the expression "z < 0.5"
z < 0.5

## [1] TRUE TRUE FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE FALSE
## [12] FALSE

## R identifies "True" with "1" and "False" with "0":
as.numeric(z < 0.5)

## [1] 1 1 0 1 0 1 0 1 1 1 0 0


## Now let x be a vector of 100 random draws from a "Normal" distribution:
x <- rnorm(100) # Generates 100 random normal observations, mean 0 sd 1

x


## [1] -0.200012990 0.699899499 -0.634473933 0.416154334 0.433994724
## [6] 0.445430480 2.996794402 0.273377078 0.357456419 0.249308401
## [11] -0.464359559 1.003507969 1.927531225 0.121692878 -1.100580925
## [16] -0.986996693 -2.139647225 0.311005285 -1.723198725 -1.084576056
## [21] -1.600534489 -0.996953194 1.049249868 1.042735406 0.675760173
## [26] 0.807134911 -1.226535930 -0.329891068 2.074956150 -1.514992320
## [31] -1.247178265 0.754553589 -0.024590894 -0.426420647 0.422274596
## [36] 2.507436861 0.379089438 -0.791699825 -1.563190554 0.971066600
## [41] -0.297439399 0.492590133 0.172821605 -1.024842560 0.609668363
## [46] -0.011182491 0.790143459 0.830592363 -0.791104395 0.631066770
## [51] 0.253769031 1.119265811 -0.089011832 -0.908859200 0.339023467
## [56] 0.948921148 0.543670492 1.685021014 -0.428554245 0.724072849
## [61] -1.885117821 1.981072195 0.873651340 0.340941382 -0.639624283
## [66] 0.283123119 -0.793733786 -1.676102155 -0.239291759 1.255400768
## [71] 1.189128889 -0.839867776 0.460160596 1.235815951 1.880859914
## [76] 0.041800470 0.001412153 -0.348594475 -1.696409877 0.758118115
## [81] 1.098939376 0.648005729 0.947246546 0.180503543 -0.853632894
## [86] -1.920442824 -1.623587718 0.135577859 -0.620126497 0.084125552
## [91] 0.064757073 0.448736428 -1.206683863 0.659806255 -1.403913311
## [96] -1.280328355 -1.341119429 -0.732032531 -1.677636407 1.299658320

par(mfrow=c(1,1)) #Resets the graphics to one plot per page
hist(x, main= "100 Obs. Standard Normal Distribution",
      breaks = 10, col = "gold") # Create a histogram of the vector x
```

100 Obs. Standard Normal Distribution




```
y <- rnorm(100) # save 100 random sample from a standard normal distribution to y  
plot(x,y)
```



```
plot(x,y,main = "Noise", col = "steelblue", pch = 18)  
abline(h=0)  
legend("topleft", title = "I'm Normal!", "Obs", pch= 18,
```


```
col = "steelblue")
```

Noise



```
plot(x,y,type="l", main = "Nonsense", col = "coral")
```

Nonsense



```
##plot
```

```
# Sampling uniformly at random, with replacement
```

```
v <- sample(1:10,100,replace=T) # Samples from 1 to 10, 100 times
```

```
v
```

```
## [1] 3 7 3 3 2 5 4 10 2 7 9 4 6 9 5 6 8 10 3 10 1 2 2
## [24] 4 10 10 5 3 7 4 4 1 7 7 4 1 9 10 5 6 10 9 5 10 9 6
## [47] 2 10 5 1 10 10 1 1 7 6 7 6 9 10 4 9 5 9 4 6 6 6 9
## [70] 8 2 7 10 1 8 5 9 9 4 6 10 2 1 5 1 6 6 5 4 2 9 6
## [93] 6 10 10 9 6 9 1 9
```

```
table(v)
```

```
## v
```

```
## 1 2 3 4 5 6 7 8 9 10
```

```
## 10 8 5 10 10 15 8 3 15 16
```

```
## Functions:
```

```
## Numerical calculations for birthday problem:
```

```
k <- 40
```

```
top <- seq(365,length=k,by=-1) # Creates a vector of 365 to 365-k
```

```
bottom <- rep(365,k) # Creates a vector filled with 365 repeated k times
```

```
top
```

```
## [1] 365 364 363 362 361 360 359 358 357 356 355 354 353 352 351 350 349
```

```
## [18] 348 347 346 345 344 343 342 341 340 339 338 337 336 335 334 333 332
```

```
## [35] 331 330 329 328 327 326
```

```
bottom
```

```
## [1] 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365
```

```
## [18] 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365 365
```

```
## [35] 365 365 365 365 365 365
```

```
top/bottom
```

```
## [1] 1.0000000 0.9972603 0.9945205 0.9917808 0.9890411 0.9863014 0.9835616
```

```
## [8] 0.9808219 0.9780822 0.9753425 0.9726027 0.9698630 0.9671233 0.9643836
```

```
## [15] 0.9616438 0.9589041 0.9561644 0.9534247 0.9506849 0.9479452 0.9452055
```

```
## [22] 0.9424658 0.9397260 0.9369863 0.9342466 0.9315068 0.9287671 0.9260274
```

```
## [29] 0.9232877 0.9205479 0.9178082 0.9150685 0.9123288 0.9095890 0.9068493
```

```
## [36] 0.9041096 0.9013699 0.8986301 0.8958904 0.8931507
```



```
prod(top/bottom) # This is the prob of NO birthday match
```

```
## [1] 0.1087682
```

```
1 - prod(top/bottom) # This is the prob of having a birthday match
```

```
## [1] 0.8912318
```

```
## Let's make a function out of what we just did:
```

```
bday <- function(k){ # k is the variable  
  top <- seq(365,length=k,by=-1)  
  bottom <- rep(365,k)  
  return(1-prod(top/bottom))  
}
```

```
bday(40)
```

```
## [1] 0.8912318
```

```
#####  
## Intro to for loops in R:  
#####
```

```
s <- 0  
for(i in 1:100){  
  s <- s+i  
}  
s
```

```
## [1] 5050
```

```
## Sometimes you can do the same thing without a loop:  
sum(1:100)
```

```
## [1] 5050
```

```
## You can have more commands in the body of the loop:
```

```
s <- 0  
for(i in 1:10){  
  s <- s+i  
  cat("When i = ", i, ", s = ", s, "\n", sep="") # "cat" prints things  
}
```

```
## When i = 1, s = 1  
## When i = 2, s = 3  
## When i = 3, s = 6  
## When i = 4, s = 10  
## When i = 5, s = 15  
## When i = 6, s = 21
```

```
## When i = 7, s = 28
## When i = 8, s = 36
## When i = 9, s = 45
## When i = 10, s = 55
```

```
s <- 0
for(i in 1:10){
  s <- s+i
  cat("When i = ", i, ", s = ", s, "\n", sep="")
  remainder2 <- (i %% 2)
  twos <- i/2
  if(remainder2 == 0){
    cat("I'm getting", rep("really", twos), "tired!\n")
  }
}
```

```
## When i = 1, s = 1
## When i = 2, s = 3
## I'm getting really tired!
## When i = 3, s = 6
## When i = 4, s = 10
## I'm getting really really tired!
## When i = 5, s = 15
## When i = 6, s = 21
## I'm getting really really really tired!
## When i = 7, s = 28
## When i = 8, s = 36
## I'm getting really really really really tired!
## When i = 9, s = 45
## When i = 10, s = 55
## I'm getting really really really really really tired!
```